

Monitoring Latencies

How fast is your REST service?

Fabian Stäber

Devoxx Belgium 2022



Fabian Stäber



Engineering Manager at Grafana Labs



Prometheus team member; maintainer of the Prometheus Java client library



@fstabr



Contents

Part 1
(demo)

```
long start = System.nanoTime();
```

```
try {
```

```
    // do something
```

```
} finally {
```

```
    long duration = System.nanoTime() - start;
```

```
    observe(duration);
```

Part 2
(slides)

```
}
```



Contents

Part 1
(demo)

- Response time vs service time
- Load test backs off when server is at its limit

Part 2
(slides)

- Average & Maximum
- Percentiles
- Histograms
- More Histograms



Evaluation Criteria

- **Time Interval**
 - Ad-hoc at query time, or predefined in the instrumentation library?
- **Aggregation**
 - Across multiple instances of a service?
- **Practical Value**
 - Opinionated



Average

Total time serving requests ←

```
increase(http_server_requests_seconds_sum[5m]) /  
increase(http_server_requests_seconds_count[5m])
```

Total number of requests ←

You cannot calculate an average of averages.
But you can sum up total times and total counts.



Average

- **Time Interval**
 - Ad-hoc at query time, like [5m] (*)
- **Aggregation**
 - Yes (*)
- **Practical Value**
 - Available by default in Spring Boot (Micrometer)
 - Affected by outliers. Example: 1000 calls take 20ms each; 1 client hangs and times out after 30s → average = 50ms

(*) if the instrumentation provides sum and count individually.



Maximum

`http_server_requests_seconds_max`

↳ As provided by Micrometer

`max(http_server_requests_seconds_max)`

↳ Aggregation



Maximum

- **Time Interval**

- Predefined in the instrumentation library
- Spring Boot (Micrometer): Sliding window, default 3 minutes, moved forward every 1 minute

- **Aggregation**

- Yes


- **Practical Value**

- Good for load tests
- Limited value for production:
 - If you are on-call, do you want to be paged when a single outlier takes long?
 - Outliers often caused by client issues, nothing's wrong with your REST service.

Recommended talk: "How NOT to Measure Latency", Gil Tene, 2015



Percentiles



Prometheus client_java does not directly support max,
but you can use the 100th percentile for max and the 0th percentile for min.

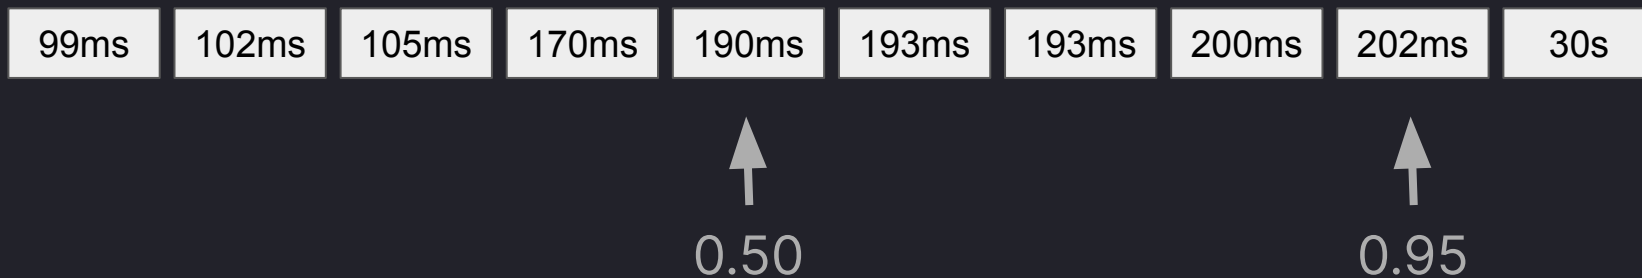


Percentiles



Expensive:

Sorted list with all observations.



Percentiles



CKMS algorithm:

Allow some error to save space.

```
Summary.build()  
  .quantile(0.5, 0.01)  
  .quantile(0.95, 0.005)  
  .quantile(0.99, 0.001);
```



G. Cormode, F. Korn, S. Muthukrishnan and D. Srivastava.
Effective Computation of Biased Quantiles over Data Streams



Percentiles



	<code>.quantile(0.95, 0.005)</code>	<code>.quantile(0.95, 0.005)</code> <code>.quantile(0.99, 0.001)</code>	<code>.quantile(0.5, 0.01)</code> <code>.quantile(0.95, 0.005)</code> <code>.quantile(0.99, 0.001)</code>
1,000	139	152	200
10,000	43	56	90
100,000	60	70	103
1,000,000	92	102	132
10,000,000	27	40	68
100,000,000	27	36	67

<https://grafana.com/blog/2022/03/01/how-summary-metrics-work-in-prometheus/>



Percentiles




- **Time Interval**
 - Predefined in the instrumentation library
 - Prometheus: Sliding window, default 10 minutes, moved forward every 2 minutes
- **Aggregation**
 - No
- **Practical Value**
 - Percentiles are awesome, also good for alerting
 - Lack of aggregation is a bummer

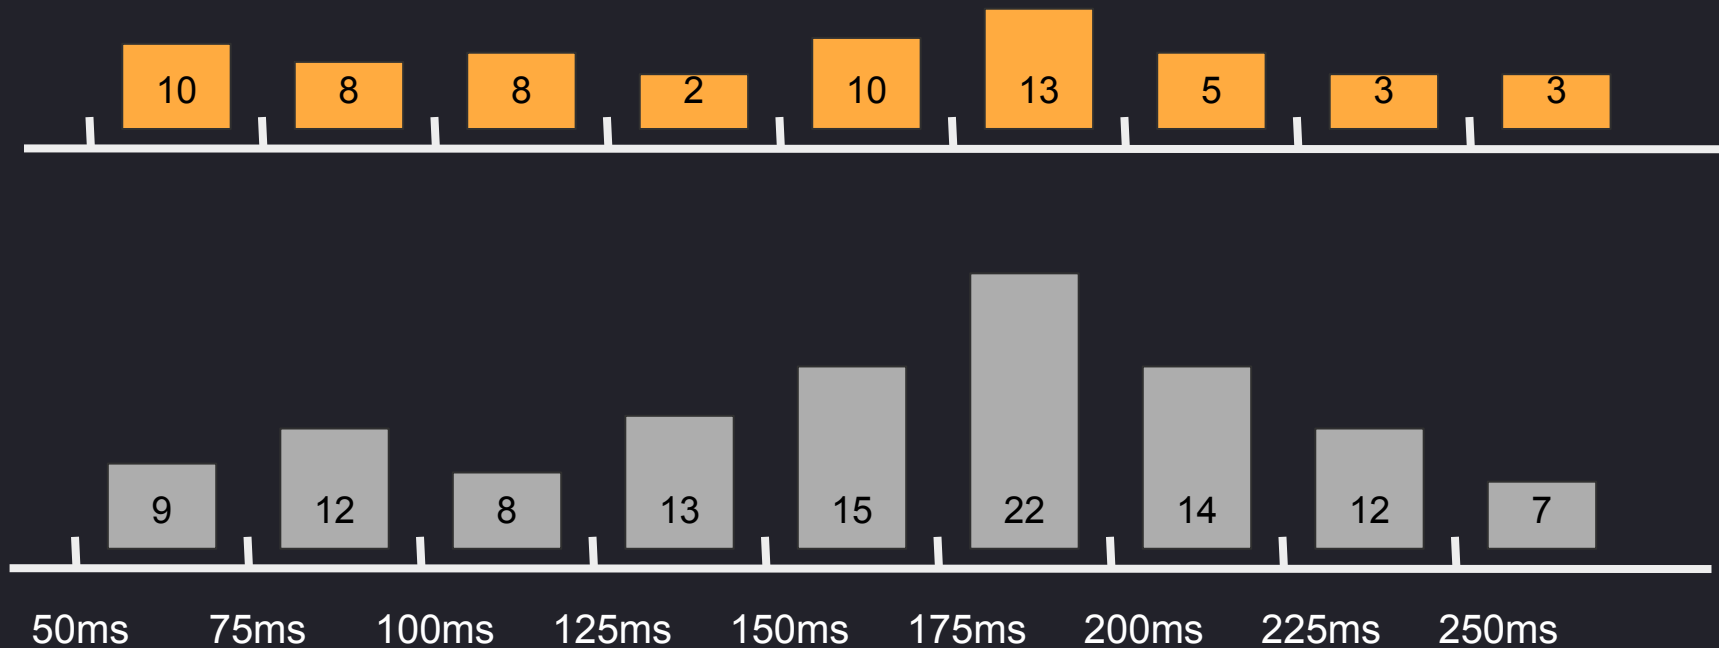
**You cannot calculate an average of averages.
You cannot calculate an average of percentiles.**




Histograms



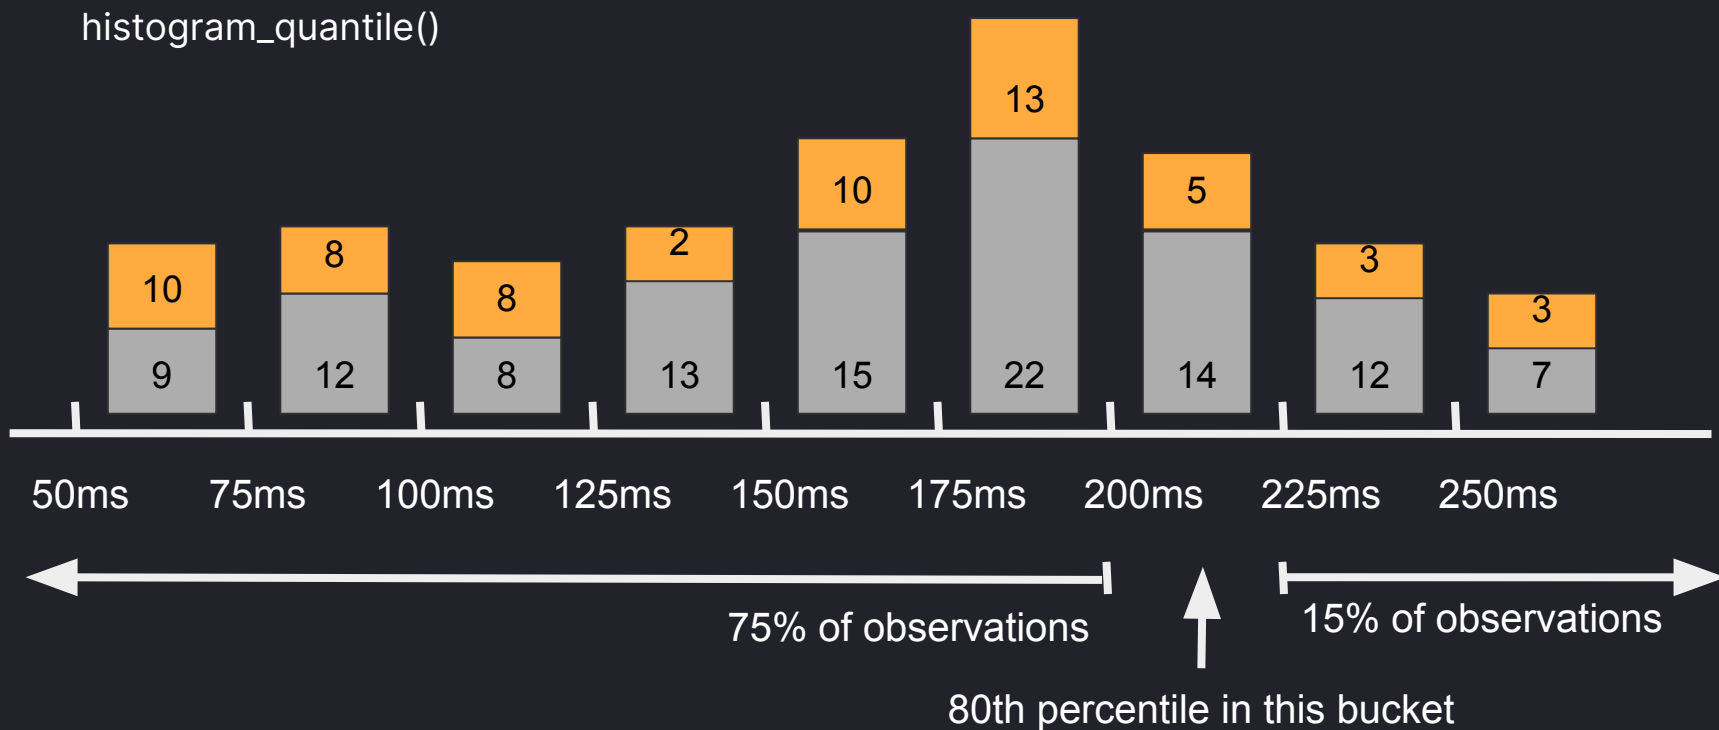
Histograms




Histograms



Histograms



Histograms



`rate(http_request_duration_seconds_bucket[5m])`




Histograms



`histogram_quantile(0.95, rate(http_request_duration_seconds_bucket[5m]))`



Histograms



`rate(http_request_duration_seconds_bucket[5m])`



Histograms



- **Time Interval**
 - Ad-hoc at query time, like [5m]
- **Aggregation**
 - Yes
- **Practical Value**
 - Awesome, but ...
 - ... you have to define reasonable bucket boundaries.



Sparse / Native / Exponential Histograms


schema=0



Sparse / Native / Exponential Histograms



Sparse / Native / Exponential Histograms




Sparse / Native / Exponential Histograms

schema	Max percentile error relative to the observed value (assuming harmonic mean)
0	33%
1	17%
2	9%
3	4%
4	2%
5	1%
6	0.5%
7	0.3%
8	0.1%



Sparse / Native / Exponential Histograms



Configure with max number of buckets:

→ If max number is reached, schema is reduced, two buckets collapse into one.



Sparse / Native / Exponential Histograms

- **Time Interval**
 - Ad-hoc at query time, like [5m]
- **Aggregation**
 - Yes
- **Practical Value**
 - No need for explicit bucket boundaries.
 - Can derive percentiles with guaranteed error limit.
 - It's new. There's little practical experience yet.



Contents

Part 1
(demo)

- Response time vs service time
- Load test backs off when server is at its limit

Part 2
(slides)

- Average & Maximum
- Percentiles
- Histograms
- More Histograms



Reference: Java Metric Libraries



- Dropwizard Metrics



- Micrometer (Spring Boot, Quarkus, ...)



- Prometheus Java client library



- OpenTelemetry Java SDK



- Microprofile Metrics 5.0



Dropwizard Metrics



- **Average**
 - Pre-calculated averages, no separate sum and count.
- **Maximum**
 - Yes, as part of Histogram
- **Percentiles**
 - Yes, as part of Histogram
- **Histograms**
 - No. There is a data type named histogram, but it exposes only percentiles, no histogram buckets.
- **Sparse / Native / Exponential Histograms**
 - No



Micrometer



- **Average**
 - Yes, using sum and count as shown in this presentation.
- **Maximum**
 - Yes
- **Percentiles**
 - Yes, using HdrHistograms internally.
- **Histograms**
 - Yes
- **Sparse / Native / Exponential Histograms**
 - No, but I assume Micrometer will add support if these histograms become popular in Prometheus



Prometheus Java Client Library



- Average
 - Yes, using sum and count (Summary in Prometheus terminology)
- Maximum
 - Yes, as the 100th percentile (of a Summary metric)
- Percentiles
 - Yes, using the CKMS algorithm described in this talk
- Histograms
 - Yes
- Sparse / Native / Exponential Histograms
 - Yes, currently in the sparsehistogram feature branch.



OpenTelemetry Metrics Java SDK



- **Average**
 - Yes, count and sum are part of Histogram
- **Maximum**
 - Yes, part of Histogram
- **Percentiles**
 - No (?)
- **Histograms**
 - Yes
- **Sparse / Native / Exponential Histograms**
 - Yes, but exposition in Prometheus format still a TODO.



Microprofile Metrics 5.0



- **Average**
 - Yes, count and sum are part of timers
- **Maximum**
 - Yes, part of timers
- **Percentiles**
 - Yes
- **Histograms**
 - Not in 5.0, planned for 5.1
- **Sparse / Native / Exponential Histograms**
 - No

